

# jjoy6\_Assignment\_4

April 20, 2025

## Assignment\_4

```
[ ]: import os, pathlib, shutil, random
      from tensorflow import keras
      from tensorflow.keras import layers
      from tensorflow.keras.layers import TextVectorization
      import tensorflow as tf
```

## Downloading IMBD data

```
[ ]: !curl -O https://ai.stanford.edu/~amaas/data/sentiment/aclImdb_v1.tar.gz
      !tar -xf aclImdb_v1.tar.gz
      !rm -r aclImdb/train/unsup
```

% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Current
			Dload Upload	Total	Spent	Left	Speed
100 80.2M	100 80.2M	0 0	18.0M	0	0:00:04	0:00:04	--:--:-- 18.0M

## Preparing Data

```
[ ]: batch_size = 32
      base_dir = pathlib.Path("aclImdb")
      val_dir = base_dir / "val"
      train_dir = base_dir / "train"
      for category in ("neg", "pos"):
          os.makedirs(val_dir / category, exist_ok=True)
          files = os.listdir(train_dir / category)
          random.Random(1337).shuffle(files)
          num_val_samples = 5000
          val_files = files[-num_val_samples:]
          for fname in val_files:
              shutil.move(train_dir / category / fname,
                          val_dir / category / fname)
```

## Function to change training size

```
[ ]: def make_subset(subset_name, new_subset_name, sample_size):
      for category in ("neg", "pos"):
          dir = base_dir / new_subset_name / category
```

```

src_dir = base_dir / subset_name / category
os.makedirs(dir, exist_ok=True)
fnames = [f for f in os.listdir(src_dir)]

#sampled_files = random.sample(fnames, sample_size)
sampled_files = fnames[:sample_size]
for sampled_files in sampled_files:
    shutil.copyfile(src=src_dir / sampled_files,
                    dst=dir / sampled_files)

```

Making subsets of various training sizes

```

[ ]: """
make_subset("train", "train100", 50)
train_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/train100", batch_size=batch_size
)
"""

"""
make_subset("train", "train200", 100)
train_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/train200", batch_size=batch_size
)
"""

"""
make_subset("train", "train300", 150)
train_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/train300", batch_size=batch_size
)
"""

"""
make_subset("train", "train7500", 7500)
train_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/train7500", batch_size=batch_size
)
"""

make_subset("train", "train7500", 25)
train_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/train7500", batch_size=batch_size
)

```

Found 50 files belonging to 2 classes.

Splitting Data into training, validation, and testing

```
[ ]: val_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/val", batch_size=batch_size
)
test_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/test", batch_size=batch_size
)
text_only_train_ds = train_ds.map(lambda x, y: x)
```

Found 10000 files belonging to 2 classes.

Found 25000 files belonging to 2 classes.

Preparing text data Text standardization Text splitting (tokenization) Vocabulary indexing Using the TextVectorization layer

```
[ ]: #Two approaches for representing groups of words: Sets and sequences
#Preparing the IMDB movie reviews data

#Processing words as a set: The bag-of-words approach
#Single words (unigrams) with binary encoding
#Preprocessing our datasets with a TextVectorization layer
tokenSize = 2771

text_vectorization = TextVectorization(
    max_tokens=10000,
    output_mode="multi_hot",
)

text_only_train_ds = train_ds.map(lambda x, y: x)
text_vectorization.adapt(text_only_train_ds)

binary_1gram_train_ds = train_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
binary_1gram_val_ds = val_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
binary_1gram_test_ds = test_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)

#Inspecting the output of our binary unigram dataset

for inputs, targets in binary_1gram_train_ds:
    print("inputs.shape:", inputs.shape)
    print("inputs.dtype:", inputs.dtype)
    print("targets.shape:", targets.shape)
    print("targets.dtype:", targets.dtype)
    print("inputs[0]:", inputs[0])
```

```

print("targets[0]:", targets[0])
break

#Our model-building utility

#With so little training data I do not have enough for 10000 tokens?

def get_model(max_tokens=tokenSize, hidden_dim=16):
    inputs = keras.Input(shape=(max_tokens,))
    x = layers.Dense(hidden_dim, activation="relu")(inputs)
    x = layers.Dropout(0.5)(x)
    outputs = layers.Dense(1, activation="sigmoid")(x)
    model = keras.Model(inputs, outputs)
    model.compile(optimizer="rmsprop",
                  loss="binary_crossentropy",
                  metrics=["accuracy"])
    return model

#Training and testing the binary unigram model

model = get_model()
model.summary()
callbacks = [
    keras.callbacks.ModelCheckpoint("binary_1gram{epoch:02d}.keras",
                                    save_best_only=True)
]

model.fit(binary_1gram_train_ds.cache(),
          validation_data=binary_1gram_val_ds.cache(),
          epochs=10,
          callbacks=callbacks)

```

```

inputs.shape: (32, 2771)
inputs.dtype: <dtype: 'int64'>
targets.shape: (32,)
targets.dtype: <dtype: 'int32'>
inputs[0]: tf.Tensor([0 1 1 ... 0 0 0], shape=(2771,), dtype=int64)
targets[0]: tf.Tensor(1, shape=(), dtype=int32)

Model: "functional_1"

```

Layer (type)	Output Shape	Param #
input_layer_1 ( <a href="#">InputLayer</a> )	( <a href="#">None</a> , 2771)	0
dense_2 ( <a href="#">Dense</a> )	( <a href="#">None</a> , 16)	44,352

dropout_1 (Dropout)	(None, 16)	0
dense_3 (Dense)	(None, 1)	17

Total params: 44,369 (173.32 KB)

Trainable params: 44,369 (173.32 KB)

Non-trainable params: 0 (0.00 B)

```
Epoch 1/10
2/2          4s 2s/step -
accuracy: 0.7433 - loss: 0.6375 - val_accuracy: 0.5227 - val_loss: 0.6907
Epoch 2/10
2/2          1s 643ms/step -
accuracy: 0.7462 - loss: 0.5791 - val_accuracy: 0.5419 - val_loss: 0.6839
Epoch 3/10
2/2          1s 621ms/step -
accuracy: 0.8175 - loss: 0.4892 - val_accuracy: 0.5405 - val_loss: 0.6845
Epoch 4/10
2/2          1s 649ms/step -
accuracy: 0.8175 - loss: 0.4740 - val_accuracy: 0.5811 - val_loss: 0.6715
Epoch 5/10
2/2          1s 610ms/step -
accuracy: 0.9525 - loss: 0.3956 - val_accuracy: 0.5797 - val_loss: 0.6691
Epoch 6/10
2/2          1s 602ms/step -
accuracy: 0.9629 - loss: 0.3485 - val_accuracy: 0.5790 - val_loss: 0.6695
Epoch 7/10
2/2          1s 622ms/step -
accuracy: 0.9525 - loss: 0.3240 - val_accuracy: 0.5884 - val_loss: 0.6650
Epoch 8/10
2/2          1s 629ms/step -
accuracy: 0.9629 - loss: 0.3042 - val_accuracy: 0.6031 - val_loss: 0.6587
Epoch 9/10
2/2          1s 656ms/step -
accuracy: 1.0000 - loss: 0.2517 - val_accuracy: 0.6110 - val_loss: 0.6561
Epoch 10/10
2/2          1s 631ms/step -
accuracy: 0.9762 - loss: 0.2649 - val_accuracy: 0.6196 - val_loss: 0.6520
```

[ ]: <keras.src.callbacks.history.History at 0x79a620b4fed0>

Testing binary\_lgram

```
[ ]: model = keras.models.load_model("binary_1gram10.keras")
print(f"Test acc: {model.evaluate(binary_1gram_test_ds)[1]:.3f}")
```

```
782/782          2s 2ms/step -
accuracy: 0.6212 - loss: 0.6524
Test acc: 0.622
```

Bigrams

```
[ ]: #Bigrams with binary encoding
#Configuring the TextVectorization layer to return bigrams

text_vectorization = TextVectorization(
    ngrams=2,
    max_tokens=tokenSize,
    output_mode="multi_hot",
)

#Training and testing the binary bigram model

text_vectorization.adapt(text_only_train_ds)
binary_2gram_train_ds = train_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
binary_2gram_val_ds = val_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
binary_2gram_test_ds = test_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)

model = get_model()
model.summary()
callbacks = [
    keras.callbacks.ModelCheckpoint("binary_2gram{epoch:02d}.keras",
                                    save_best_only=True)
]
model.fit(binary_2gram_train_ds.cache(),
          validation_data=binary_2gram_val_ds.cache(),
          epochs=10,
          callbacks=callbacks)
```

Model: "functional\_2"

Layer (type)	Output Shape	Param #
input_layer_2 ( <a href="#">InputLayer</a> )	( <a href="#">None</a> , 2771)	0

dense_4 (Dense)	(None, 16)	44,352
dropout_2 (Dropout)	(None, 16)	0
dense_5 (Dense)	(None, 1)	17

Total params: 44,369 (173.32 KB)

Trainable params: 44,369 (173.32 KB)

Non-trainable params: 0 (0.00 B)

Epoch 1/10

2/2 3s 2s/step -  
accuracy: 0.4288 - loss: 0.7595 - val\_accuracy: 0.5139 - val\_loss: 0.6963

Epoch 2/10

2/2 1s 625ms/step -  
accuracy: 0.7167 - loss: 0.6158 - val\_accuracy: 0.5357 - val\_loss: 0.6887

Epoch 3/10

2/2 1s 634ms/step -  
accuracy: 0.8071 - loss: 0.4952 - val\_accuracy: 0.5422 - val\_loss: 0.6864

Epoch 4/10

2/2 1s 634ms/step -  
accuracy: 0.9287 - loss: 0.4509 - val\_accuracy: 0.5653 - val\_loss: 0.6785

Epoch 5/10

2/2 1s 634ms/step -  
accuracy: 0.9525 - loss: 0.3550 - val\_accuracy: 0.5754 - val\_loss: 0.6746

Epoch 6/10

2/2 1s 614ms/step -  
accuracy: 1.0000 - loss: 0.2999 - val\_accuracy: 0.5740 - val\_loss: 0.6750

Epoch 7/10

2/2 1s 642ms/step -  
accuracy: 0.9021 - loss: 0.3184 - val\_accuracy: 0.5998 - val\_loss: 0.6656

Epoch 8/10

2/2 1s 637ms/step -  
accuracy: 0.9287 - loss: 0.2567 - val\_accuracy: 0.6067 - val\_loss: 0.6620

Epoch 9/10

2/2 1s 635ms/step -  
accuracy: 0.9525 - loss: 0.2514 - val\_accuracy: 0.6130 - val\_loss: 0.6538

Epoch 10/10

2/2 1s 618ms/step -  
accuracy: 1.0000 - loss: 0.1964 - val\_accuracy: 0.6183 - val\_loss: 0.6545

```
[ ]: <keras.src.callbacks.history.History at 0x79a61c2daf50>
```

Testing Bigram

```
[ ]: model = keras.models.load_model("binary_2gram09.keras")
print(f"Test acc: {model.evaluate(binary_2gram_test_ds)[1]:.3f}")
```

```
782/782          2s 2ms/step -
accuracy: 0.6206 - loss: 0.6535
Test acc: 0.618
```

Bigram using TF-IDF encoding

```
[ ]: #Configuring TextVectorization to return TF-IDF-weighted outputs

text_vectorization = TextVectorization(
    ngrams=2,
    max_tokens=tokenSize,
    output_mode="tf_idf",
)

#Training and testing the TF-IDF bigram model

text_vectorization.adapt(text_only_train_ds)

tfidf_2gram_train_ds = train_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
tfidf_2gram_val_ds = val_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
tfidf_2gram_test_ds = test_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)

model = get_model()
model.summary()

callbacks = [
    keras.callbacks.ModelCheckpoint("tfidf_2gram{epoch:02d}.keras",
                                    save_best_only=True)
]
model.fit(tfidf_2gram_train_ds.cache(),
          validation_data=tfidf_2gram_val_ds.cache(),
          epochs=10,
          callbacks=callbacks)
```

Model: "functional\_3"



Layer (type)	Output Shape	Param #
input_layer_3 (InputLayer)	(None, 2771)	0
dense_6 (Dense)	(None, 16)	44,352
dropout_3 (Dropout)	(None, 16)	0
dense_7 (Dense)	(None, 1)	17

Total params: 44,369 (173.32 KB)

Trainable params: 44,369 (173.32 KB)

Non-trainable params: 0 (0.00 B)

Epoch 1/10

2/2 3s 2s/step -

accuracy: 0.5104 - loss: 8.9538 - val\_accuracy: 0.5000 - val\_loss: 6.9353

Epoch 2/10

2/2 1s 608ms/step -

accuracy: 0.5504 - loss: 4.0754 - val\_accuracy: 0.5000 - val\_loss: 4.8370

Epoch 3/10

2/2 1s 601ms/step -

accuracy: 0.5446 - loss: 4.3268 - val\_accuracy: 0.5001 - val\_loss: 3.3717

Epoch 4/10

2/2 1s 609ms/step -

accuracy: 0.6662 - loss: 1.6967 - val\_accuracy: 0.5003 - val\_loss: 2.7141

Epoch 5/10

2/2 1s 610ms/step -

accuracy: 0.5504 - loss: 2.3754 - val\_accuracy: 0.5004 - val\_loss: 2.5636

Epoch 6/10

2/2 1s 605ms/step -

accuracy: 0.7404 - loss: 0.7364 - val\_accuracy: 0.5012 - val\_loss: 2.2393

Epoch 7/10

2/2 1s 599ms/step -

accuracy: 0.6542 - loss: 1.1727 - val\_accuracy: 0.5029 - val\_loss: 1.9939

Epoch 8/10

2/2 1s 604ms/step -

accuracy: 0.7508 - loss: 0.5280 - val\_accuracy: 0.5057 - val\_loss: 1.8550

Epoch 9/10

2/2 1s 590ms/step -

accuracy: 0.7937 - loss: 0.5456 - val\_accuracy: 0.5038 - val\_loss: 1.9760

Epoch 10/10

```
2/2          1s 591ms/step -
accuracy: 0.8442 - loss: 0.4268 - val_accuracy: 0.5028 - val_loss: 2.2357
```

```
[ ]: <keras.src.callbacks.history.History at 0x79a5fea17bd0>
```

Testing Bigram with TF-IDF

```
[ ]: model = keras.models.load_model("tfidf_2gram08.keras")
print(f"Test acc: {model.evaluate(tfidf_2gram_test_ds)[1]:.3f}")
```

```
782/782          2s 2ms/step -
accuracy: 0.5039 - loss: 1.7967
Test acc: 0.505
```

Sequence model with one hot encoding

```
[ ]: #Preparing integer sequence datasets

max_length = 150
max_tokens = 10000
text_vectorization = layers.TextVectorization(
    max_tokens=max_tokens,
    output_mode="int",
    output_sequence_length=max_length,
)

text_vectorization.adapt(text_only_train_ds)

int_train_ds = train_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
int_val_ds = val_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
int_test_ds = test_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)

#A sequence model built on one-hot encoded vector sequences

inputs = keras.Input(shape=(None,), dtype="int64")
#embedded = tf.keras.layers.Lambda(lambda x: tf.one_hot(x, depth=max_tokens),  
↪output_shape=(None, max_tokens))(inputs)
embedded = keras.ops.one_hot(inputs, num_classes=max_tokens)
x = layers.Bidirectional(layers.LSTM(32))(embedded)
x = layers.Dropout(0.5)(x)
```

```

outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
model.summary()

#Training a first basic sequence model

callbacks = [
    keras.callbacks.ModelCheckpoint("one_hot_bidir_lstm{epoch:02d}.keras",
                                   save_best_only=True)
]
model.fit(int_train_ds, validation_data=int_val_ds, epochs=10,
        ↪callbacks=callbacks)

#model.evaluate(int_test_ds)

```

Model: "functional\_4"

Layer (type)	Output Shape	Param #
input_layer_4 ( <a href="#">InputLayer</a> )	(None, None)	0
one_hot ( <a href="#">OneHot</a> )	(None, None, 10000)	0
bidirectional ( <a href="#">Bidirectional</a> )	(None, 64)	2,568,448
dropout_4 ( <a href="#">Dropout</a> )	(None, 64)	0
dense_8 ( <a href="#">Dense</a> )	(None, 1)	65

Total params: 2,568,513 (9.80 MB)

Trainable params: 2,568,513 (9.80 MB)

Non-trainable params: 0 (0.00 B)

Epoch 1/10

2/2                    6s 3s/step -

accuracy: 0.4838 - loss: 0.6939 - val\_accuracy: 0.5077 - val\_loss: 0.6930

Epoch 2/10

```

2/2          3s 3s/step -
accuracy: 0.6721 - loss: 0.6911 - val_accuracy: 0.5150 - val_loss: 0.6929
Epoch 3/10
2/2          3s 3s/step -
accuracy: 0.8042 - loss: 0.6866 - val_accuracy: 0.5178 - val_loss: 0.6929
Epoch 4/10
2/2          3s 3s/step -
accuracy: 0.7596 - loss: 0.6836 - val_accuracy: 0.5203 - val_loss: 0.6927
Epoch 5/10
2/2          3s 3s/step -
accuracy: 0.8754 - loss: 0.6809 - val_accuracy: 0.5182 - val_loss: 0.6927
Epoch 6/10
2/2          3s 3s/step -
accuracy: 0.8546 - loss: 0.6758 - val_accuracy: 0.5188 - val_loss: 0.6926
Epoch 7/10
2/2          3s 3s/step -
accuracy: 0.9392 - loss: 0.6705 - val_accuracy: 0.5203 - val_loss: 0.6924
Epoch 8/10
2/2          3s 3s/step -
accuracy: 0.9125 - loss: 0.6665 - val_accuracy: 0.5213 - val_loss: 0.6922
Epoch 9/10
2/2          3s 3s/step -
accuracy: 0.9021 - loss: 0.6633 - val_accuracy: 0.5245 - val_loss: 0.6920
Epoch 10/10
2/2          3s 3s/step -
accuracy: 0.9258 - loss: 0.6540 - val_accuracy: 0.5293 - val_loss: 0.6917

```

```
[ ]: <keras.src.callbacks.history.History at 0x79a5d39a9ed0>
```

Testing Sequence with one hot encoding

```
[ ]: model = keras.models.load_model("one_hot_bidir_lstm10.keras", safe_mode=False)

print(f"Test acc: {model.evaluate(int_test_ds)[1]:.3f}")
```

```

782/782          7s 8ms/step -
accuracy: 0.5265 - loss: 0.6919
Test acc: 0.526

```

Embedding Layer from scratch

```
[ ]: #Understanding word embeddings
#Learning word embeddings with the Embedding layer
#Instantiating an Embedding layer

embedding_layer = layers.Embedding(input_dim=max_tokens, output_dim=256)

#Model that uses an Embedding layer trained from scratch
```

```

inputs = keras.Input(shape=(None,), dtype="int64")
embedded = layers.Embedding(input_dim=max_tokens, output_dim=256)(inputs)
x = layers.Bidirectional(layers.LSTM(32))(embedded)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
model.summary()

callbacks = [
    keras.callbacks.ModelCheckpoint("embeddings_bidir_gru{epoch:02d}.keras",
                                   save_best_only=True)
]
model.fit(int_train_ds, validation_data=int_val_ds, epochs=10,
        ↪callbacks=callbacks)

```

Model: "functional\_5"

Layer (type)	Output Shape	Param #
input_layer_5 ( <a href="#">InputLayer</a> )	(None, None)	0
embedding_1 ( <a href="#">Embedding</a> )	(None, None, 256)	2,560,000
bidirectional_1 ( <a href="#">Bidirectional</a> )	(None, 64)	73,984
dropout_5 ( <a href="#">Dropout</a> )	(None, 64)	0
dense_9 ( <a href="#">Dense</a> )	(None, 1)	65

Total params: 2,634,049 (10.05 MB)

Trainable params: 2,634,049 (10.05 MB)

Non-trainable params: 0 (0.00 B)

Epoch 1/10

2/2                      5s 2s/step -

accuracy: 0.5846 - loss: 0.6915 - val\_accuracy: 0.5284 - val\_loss: 0.6924

Epoch 2/10

```

2/2          2s 2s/step -
accuracy: 0.6662 - loss: 0.6811 - val_accuracy: 0.5212 - val_loss: 0.6920
Epoch 3/10
2/2          2s 2s/step -
accuracy: 0.8042 - loss: 0.6602 - val_accuracy: 0.5233 - val_loss: 0.6917
Epoch 4/10
2/2          2s 2s/step -
accuracy: 0.7538 - loss: 0.6463 - val_accuracy: 0.5250 - val_loss: 0.6912
Epoch 5/10
2/2          2s 2s/step -
accuracy: 0.8012 - loss: 0.6292 - val_accuracy: 0.5247 - val_loss: 0.6904
Epoch 6/10
2/2          2s 2s/step -
accuracy: 0.9496 - loss: 0.6007 - val_accuracy: 0.5275 - val_loss: 0.6900
Epoch 7/10
2/2          2s 2s/step -
accuracy: 0.9125 - loss: 0.5760 - val_accuracy: 0.5364 - val_loss: 0.6889
Epoch 8/10
2/2          2s 2s/step -
accuracy: 0.9333 - loss: 0.5403 - val_accuracy: 0.5296 - val_loss: 0.7019
Epoch 9/10
2/2          2s 2s/step -
accuracy: 0.8650 - loss: 0.4850 - val_accuracy: 0.5588 - val_loss: 0.7010
Epoch 10/10
2/2          2s 2s/step -
accuracy: 0.8696 - loss: 0.4584 - val_accuracy: 0.5485 - val_loss: 0.6943

```

```
[ ]: <keras.src.callbacks.history.History at 0x79a5a1927ed0>
```

Testing embedding layer from scratch

```
[ ]: model = keras.models.load_model("embeddings_bidir_gru07.keras")
print(f"Test acc: {model.evaluate(int_test_ds)[1]:.3f}")
```

```

782/782          5s 6ms/step -
accuracy: 0.5337 - loss: 0.6906
Test acc: 0.532

```

Adding masking

```
[ ]: #Understanding padding and masking
#Using an Embedding layer with masking enabled

inputs = keras.Input(shape=(None,), dtype="int64")
embedded = layers.Embedding(
    input_dim=max_tokens, output_dim=256, mask_zero=True)(inputs)
x = layers.Bidirectional(layers.LSTM(32))(embedded)
x = layers.Dropout(0.5)(x)
```

```

outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
model.summary()

callbacks = [
    keras.callbacks.ModelCheckpoint("embeddings_bidir_lstm_with_masking{epoch:
    ↪02d}.keras",
                                   save_best_only=True)
]
model.fit(int_train_ds, validation_data=int_val_ds, epochs=10,
    ↪callbacks=callbacks)

```

Model: "functional\_6"

Layer (type)	Output Shape	Param #	Connected to
input_layer_6 (InputLayer)	(None, None)	0	-
embedding_2 (Embedding)	(None, None, 256)	2,560,000	input_layer_6[0]...
not_equal (NotEqual)	(None, None)	0	input_layer_6[0]...
bidirectional_2 (Bidirectional)	(None, 64)	73,984	embedding_2[0][0]... not_equal[0][0]
dropout_6 (Dropout)	(None, 64)	0	bidirectional_2[...
dense_10 (Dense)	(None, 1)	65	dropout_6[0][0]

Total params: 2,634,049 (10.05 MB)

Trainable params: 2,634,049 (10.05 MB)

Non-trainable params: 0 (0.00 B)

Epoch 1/10

2/2                      5s 3s/step -

```

accuracy: 0.4392 - loss: 0.6946 - val_accuracy: 0.5046 - val_loss: 0.6930
Epoch 2/10
2/2          2s 2s/step -
accuracy: 0.7300 - loss: 0.6804 - val_accuracy: 0.5081 - val_loss: 0.6927
Epoch 3/10
2/2          2s 2s/step -
accuracy: 0.8308 - loss: 0.6675 - val_accuracy: 0.5120 - val_loss: 0.6924
Epoch 4/10
2/2          2s 2s/step -
accuracy: 0.9762 - loss: 0.6501 - val_accuracy: 0.5122 - val_loss: 0.6922
Epoch 5/10
2/2          2s 2s/step -
accuracy: 0.9525 - loss: 0.6289 - val_accuracy: 0.5191 - val_loss: 0.6919
Epoch 6/10
2/2          2s 2s/step -
accuracy: 0.9525 - loss: 0.6128 - val_accuracy: 0.5172 - val_loss: 0.6917
Epoch 7/10
2/2          2s 2s/step -
accuracy: 0.9867 - loss: 0.5893 - val_accuracy: 0.5223 - val_loss: 0.6915
Epoch 8/10
2/2          2s 2s/step -
accuracy: 1.0000 - loss: 0.5442 - val_accuracy: 0.5241 - val_loss: 0.6914
Epoch 9/10
2/2          2s 2s/step -
accuracy: 1.0000 - loss: 0.5091 - val_accuracy: 0.5346 - val_loss: 0.6914
Epoch 10/10
2/2          2s 2s/step -
accuracy: 1.0000 - loss: 0.4572 - val_accuracy: 0.5298 - val_loss: 0.6943

```

```
[ ]: <keras.src.callbacks.history.History at 0x79a5341ffed0>
```

Testing model that uses masking and padding

```
[ ]: model = keras.models.load_model("embeddings_bidir_lstm_with_masking09.keras")
print(f"Test acc: {model.evaluate(int_test_ds)[1]:.3f}")
```

```

782/782          6s 6ms/step -
accuracy: 0.5211 - loss: 0.6927
Test acc: 0.525

```

Downloading GloVe word-embeddings

```
[ ]: #Using pretrained word embeddings
!wget http://nlp.stanford.edu/data/glove.6B.zip
!unzip -q glove.6B.zip
```

```

--2025-04-20 15:09:46-- http://nlp.stanford.edu/data/glove.6B.zip
Resolving nlp.stanford.edu (nlp.stanford.edu)... 171.64.67.140
Connecting to nlp.stanford.edu (nlp.stanford.edu)|171.64.67.140|:80...

```



```
connected.
HTTP request sent, awaiting response... 302 Found
Location: https://nlp.stanford.edu/data/glove.6B.zip [following]
--2025-04-20 15:09:46-- https://nlp.stanford.edu/data/glove.6B.zip
Connecting to nlp.stanford.edu (nlp.stanford.edu)|171.64.67.140|:443...
connected.
HTTP request sent, awaiting response... 301 Moved Permanently
Location: https://downloads.cs.stanford.edu/nlp/data/glove.6B.zip [following]
--2025-04-20 15:09:46-- https://downloads.cs.stanford.edu/nlp/data/glove.6B.zip
Resolving downloads.cs.stanford.edu (downloads.cs.stanford.edu)... 171.64.64.22
Connecting to downloads.cs.stanford.edu
(downloads.cs.stanford.edu)|171.64.64.22|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 862182613 (822M) [application/zip]
Saving to: 'glove.6B.zip'
```

```
glove.6B.zip          100%[=====>] 822.24M  5.00MB/s    in 2m 39s
```

```
2025-04-20 15:12:25 (5.18 MB/s) - 'glove.6B.zip' saved [862182613/862182613]
```

Preparing GloVe embeddings for use

```
[ ]: #Parsing the GloVe word-embeddings file

import numpy as np
path_to_glove_file = "glove.6B.100d.txt"

embeddings_index = {}
with open(path_to_glove_file) as f:
    for line in f:
        word, coefs = line.split(maxsplit=1)
        coefs = np.fromstring(coefs, "f", sep=" ")
        embeddings_index[word] = coefs

print(f"Found {len(embeddings_index)} word vectors.")

#Preparing the GloVe word-embeddings matrix

embedding_dim = 100

vocabulary = text_vectorization.get_vocabulary()
word_index = dict(zip(vocabulary, range(len(vocabulary))))

embedding_matrix = np.zeros((max_tokens, embedding_dim))
for word, i in word_index.items():
    if i < max_tokens:
        embedding_vector = embeddings_index.get(word)
```

```

        if embedding_vector is not None:
            embedding_matrix[i] = embedding_vector
embedding_layer = layers.Embedding(
    max_tokens,
    embedding_dim,
    embeddings_initializer=keras.initializers.Constant(embedding_matrix),
    trainable=False,
    mask_zero=True,
)

```

Found 400000 word vectors.

Using Pretrained embedding

```

[ ]: #Model that uses a pretrained Embedding layer

inputs = keras.Input(shape=(None,), dtype="int64")
embedded = embedding_layer(inputs)
x = layers.Bidirectional(layers.LSTM(32))(embedded)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
model.summary()

callbacks = [
    keras.callbacks.ModelCheckpoint("glove_embeddings_sequence_model{epoch:02d}.
    ↪keras",
                                save_best_only=True)
]
model.fit(int_train_ds, validation_data=int_val_ds, epochs=10,
    ↪callbacks=callbacks)

```

Model: "functional\_8"

Layer (type)	Output Shape	Param #	Connected to
input_layer_8 (InputLayer)	(None, None)	0	-
embedding_3 (Embedding)	(None, None, 100)	1,000,000	input_layer_8[0]...
not_equal_4 (NotEqual)	(None, None)	0	input_layer_8[0]...

bidirectional_4 (Bidirectional)	(None, 64)	34,048	embedding_3[1][0... not_equal_4[0][0]
dropout_8 (Dropout)	(None, 64)	0	bidirectional_4[...
dense_12 (Dense)	(None, 1)	65	dropout_8[0][0]

Total params: 1,034,113 (3.94 MB)

Trainable params: 34,113 (133.25 KB)

Non-trainable params: 1,000,000 (3.81 MB)

Epoch 1/10

2/2 5s 4s/step -

accuracy: 0.4704 - loss: 0.7567 - val\_accuracy: 0.5021 - val\_loss: 0.6971

Epoch 2/10

2/2 4s 4s/step -

accuracy: 0.5208 - loss: 0.6717 - val\_accuracy: 0.5211 - val\_loss: 0.6945

Epoch 3/10

2/2 2s 2s/step -

accuracy: 0.4971 - loss: 0.6807 - val\_accuracy: 0.5295 - val\_loss: 0.6946

Epoch 4/10

2/2 3s 3s/step -

accuracy: 0.5742 - loss: 0.6760 - val\_accuracy: 0.5379 - val\_loss: 0.6934

Epoch 5/10

2/2 2s 2s/step -

accuracy: 0.8250 - loss: 0.5773 - val\_accuracy: 0.5254 - val\_loss: 0.6940

Epoch 6/10

2/2 2s 2s/step -

accuracy: 0.6587 - loss: 0.6281 - val\_accuracy: 0.5355 - val\_loss: 0.6945

Epoch 7/10

2/2 2s 2s/step -

accuracy: 0.6217 - loss: 0.6263 - val\_accuracy: 0.5271 - val\_loss: 0.6965

Epoch 8/10

2/2 2s 2s/step -

accuracy: 0.7300 - loss: 0.5876 - val\_accuracy: 0.5322 - val\_loss: 0.6969

Epoch 9/10

2/2 2s 2s/step -

accuracy: 0.7746 - loss: 0.5828 - val\_accuracy: 0.5221 - val\_loss: 0.7026

Epoch 10/10

2/2 2s 2s/step -

accuracy: 0.7804 - loss: 0.5587 - val\_accuracy: 0.5278 - val\_loss: 0.7020

```
[ ]: <keras.src.callbacks.history.History at 0x79a5bfff79290>
```

Testing pretrained embedding

```
[ ]: model = keras.models.load_model("glove_embeddings_sequence_model04.keras")
print(f"Test acc: {model.evaluate(int_test_ds)[1]:.3f}")
```

```
782/782          6s 6ms/step -
accuracy: 0.5350 - loss: 0.6948
Test acc: 0.535
```

Transformer

```
[ ]: #Vectorizing the data

max_length = 150
max_tokens = 10000
text_vectorization = layers.TextVectorization(
    max_tokens=max_tokens,
    output_mode="int",
    output_sequence_length=max_length,
)

text_vectorization.adapt(text_only_train_ds)

int_train_ds = train_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
int_val_ds = val_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
int_test_ds = test_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
#Transformer encoder implemented as a subclassed Layer

class TransformerEncoder(layers.Layer):
    def __init__(self, embed_dim, dense_dim, num_heads, **kwargs):
        super().__init__(**kwargs)
        self.embed_dim = embed_dim
        self.dense_dim = dense_dim
        self.num_heads = num_heads
        self.attention = layers.MultiHeadAttention(
            num_heads=num_heads, key_dim=embed_dim)
        self.dense_proj = keras.Sequential(
            [layers.Dense(dense_dim, activation="relu"),
             layers.Dense(embed_dim),]
        )
```

```

self.layernorm_1 = layers.LayerNormalization()
self.layernorm_2 = layers.LayerNormalization()

def call(self, inputs, mask=None):
    if mask is not None:
        mask = mask[:, tf.newaxis, :]
    attention_output = self.attention(
        inputs, inputs, attention_mask=mask)
    proj_input = self.layernorm_1(inputs + attention_output)
    proj_output = self.dense_proj(proj_input)
    return self.layernorm_2(proj_input + proj_output)

def get_config(self):
    config = super().get_config()
    config.update({
        "embed_dim": self.embed_dim,
        "num_heads": self.num_heads,
        "dense_dim": self.dense_dim,
    })
    return config
#Using the Transformer encoder for text classification

vocab_size = 10000
embed_dim = 256
num_heads = 2
dense_dim = 32

inputs = keras.Input(shape=(None,), dtype="int64")
#x = embedding_layer(inputs)
x = layers.Embedding(vocab_size, embed_dim)(inputs)
x = TransformerEncoder(embed_dim, dense_dim, num_heads)(x)
x = layers.GlobalMaxPooling1D()(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
model.summary()
#Training and evaluating the Transformer encoder based model

callbacks = [
    keras.callbacks.ModelCheckpoint("transformer_encoder{epoch:02d}.keras",
                                   save_best_only=True)
]
model.fit(int_train_ds, validation_data=int_val_ds, epochs=20,
        ↪callbacks=callbacks)

```

Model: "functional\_8"

Layer (type)	Output Shape	Param #
input_layer_7 (InputLayer)	(None, None)	0
embedding_4 (Embedding)	(None, None, 256)	2,560,000
transformer_encoder (TransformerEncoder)	(None, None, 256)	543,776
global_max_pooling1d (GlobalMaxPooling1D)	(None, 256)	0
dropout_8 (Dropout)	(None, 256)	0
dense_12 (Dense)	(None, 1)	257

Total params: 3,104,033 (11.84 MB)

Trainable params: 3,104,033 (11.84 MB)

Non-trainable params: 0 (0.00 B)

Epoch 1/20

469/469 58s 73ms/step -

accuracy: 0.6475 - loss: 0.7964 - val\_accuracy: 0.8120 - val\_loss: 0.4147

Epoch 2/20

469/469 3s 7ms/step -

accuracy: 0.8030 - loss: 0.4326 - val\_accuracy: 0.8229 - val\_loss: 0.3852

Epoch 3/20

469/469 3s 6ms/step -

accuracy: 0.8286 - loss: 0.3797 - val\_accuracy: 0.8091 - val\_loss: 0.4239

Epoch 4/20

469/469 3s 6ms/step -

accuracy: 0.8492 - loss: 0.3507 - val\_accuracy: 0.8278 - val\_loss: 0.3877

Epoch 5/20

469/469 3s 6ms/step -

accuracy: 0.8657 - loss: 0.3165 - val\_accuracy: 0.8345 - val\_loss: 0.3720

Epoch 6/20

469/469 3s 6ms/step -

accuracy: 0.8755 - loss: 0.2929 - val\_accuracy: 0.8362 - val\_loss: 0.3863

Epoch 7/20

```

469/469          3s 6ms/step -
accuracy: 0.8917 - loss: 0.2650 - val_accuracy: 0.8420 - val_loss: 0.3838
Epoch 8/20
469/469          3s 6ms/step -
accuracy: 0.9113 - loss: 0.2274 - val_accuracy: 0.8426 - val_loss: 0.3781
Epoch 9/20
469/469          3s 6ms/step -
accuracy: 0.9202 - loss: 0.2067 - val_accuracy: 0.8419 - val_loss: 0.4218
Epoch 10/20
469/469          3s 6ms/step -
accuracy: 0.9310 - loss: 0.1814 - val_accuracy: 0.8311 - val_loss: 0.4467
Epoch 11/20
469/469          3s 6ms/step -
accuracy: 0.9416 - loss: 0.1550 - val_accuracy: 0.8303 - val_loss: 0.4882
Epoch 12/20
469/469          3s 6ms/step -
accuracy: 0.9524 - loss: 0.1311 - val_accuracy: 0.8369 - val_loss: 0.5257
Epoch 13/20
469/469          3s 6ms/step -
accuracy: 0.9631 - loss: 0.1078 - val_accuracy: 0.8255 - val_loss: 0.6154
Epoch 14/20
469/469          3s 6ms/step -
accuracy: 0.9691 - loss: 0.0911 - val_accuracy: 0.8333 - val_loss: 0.6085
Epoch 15/20
469/469          3s 6ms/step -
accuracy: 0.9724 - loss: 0.0800 - val_accuracy: 0.8277 - val_loss: 0.6675
Epoch 16/20
469/469          3s 7ms/step -
accuracy: 0.9806 - loss: 0.0577 - val_accuracy: 0.8333 - val_loss: 0.7185
Epoch 17/20
469/469          3s 6ms/step -
accuracy: 0.9818 - loss: 0.0531 - val_accuracy: 0.8314 - val_loss: 0.7806
Epoch 18/20
469/469          3s 6ms/step -
accuracy: 0.9837 - loss: 0.0485 - val_accuracy: 0.8269 - val_loss: 0.8679
Epoch 19/20
469/469          3s 6ms/step -
accuracy: 0.9876 - loss: 0.0377 - val_accuracy: 0.8244 - val_loss: 0.9400
Epoch 20/20
469/469          3s 6ms/step -
accuracy: 0.9887 - loss: 0.0319 - val_accuracy: 0.8272 - val_loss: 0.9225

```

```
[ ]: <keras.src.callbacks.history.History at 0x7adb1ff70610>
```

Testing transformer

```
[ ]: model = keras.models.load_model(
    "transformer_encoder05.keras",
    custom_objects={"TransformerEncoder": TransformerEncoder})
print(f"Test acc: {model.evaluate(int_test_ds)[1]:.3f}")
```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/layer.py:393:  
 UserWarning: `build()` was called on layer 'transformer\_encoder', however the layer does not have a `build()` method implemented and it looks like it has unbuilt state. This will cause the layer to be marked as built, despite not being actually built, which may cause failures down the line. Make sure to implement a proper `build()` method.

```
warnings.warn(
```

```
782/782          8s 9ms/step -
accuracy: 0.8304 - loss: 0.3765
Test acc: 0.830
```

Transformer with positional embedding

```
[ ]: #Using positional encoding to re-inject order information
      #Implementing positional embedding as a subclassed layer

class PositionalEmbedding(layers.Layer):
    def __init__(self, sequence_length, input_dim, output_dim, **kwargs):
        super().__init__(**kwargs)
        self.token_embeddings = layers.Embedding(
            input_dim=input_dim, output_dim=output_dim)
        self.position_embeddings = layers.Embedding(
            input_dim=sequence_length, output_dim=output_dim)
        self.sequence_length = sequence_length
        self.input_dim = input_dim
        self.output_dim = output_dim

    def call(self, inputs):
        length = tf.shape(inputs)[-1]
        positions = tf.range(start=0, limit=length, delta=1)
        embedded_tokens = self.token_embeddings(inputs)
        embedded_positions = self.position_embeddings(positions)
        return embedded_tokens + embedded_positions

    def compute_mask(self, inputs, mask=None):
        return keras.ops.not_equal(inputs, 0)

    def get_config(self):
        config = super().get_config()
        config.update({
            "output_dim": self.output_dim,
            "sequence_length": self.sequence_length,
```



```

        "input_dim": self.input_dim,
    })
    return config
#Putting it all together: A text-classification Transformer
#Combining the Transformer encoder with positional embedding

vocab_size = 10000
sequence_length = 150
embed_dim = 256
num_heads = 2
dense_dim = 32

inputs = keras.Input(shape=(None,), dtype="int64")
x = PositionalEmbedding(sequence_length, vocab_size, embed_dim)(inputs)
x = TransformerEncoder(embed_dim, dense_dim, num_heads)(x)
x = layers.GlobalMaxPooling1D()(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
model.summary()

callbacks = [
    keras.callbacks.ModelCheckpoint("full_transformer_encoder{epoch:02d}.keras",
                                   save_best_only=True)
]
model.fit(int_train_ds, validation_data=int_val_ds, epochs=20,
        ↪callbacks=callbacks)

```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/layer.py:938:  
UserWarning: Layer 'transformer\_encoder\_1' (of type TransformerEncoder) was  
passed an input with a mask attached to it. However, this layer does not support  
masking and will therefore destroy the mask information. Downstream layers will  
not see the mask.

warnings.warn(

Model: "functional\_11"

Layer (type)	Output Shape	Param #	Connected to
input_layer_10 (InputLayer)	(None, None)	0	-
positional_embeddi...	(None, None, 256)	2,598,400	input_layer_10[0...

```

(PositionalEmbeddi...

not_equal_4          (None, None)          0  input_layer_10[0...
(NotEqual)

transformer_encode... (None, None, 256)    543,776  positional_embed...
(TransformerEncode... not_equal_4[0][0]

global_max_pooling... (None, 256)          0  transformer_enco...
(GlobalMaxPooling1...

dropout_11           (None, 256)          0  global_max_pooli...
(Dropout)

dense_17 (Dense)      (None, 1)            257  dropout_11[0][0]

```

Total params: 3,142,433 (11.99 MB)

Trainable params: 3,142,433 (11.99 MB)

Non-trainable params: 0 (0.00 B)

Epoch 1/20

```

469/469          30s 38ms/step -
accuracy: 0.6043 - loss: 0.8663 - val_accuracy: 0.8066 - val_loss: 0.4306
Epoch 2/20

```

```

469/469          3s 6ms/step -
accuracy: 0.7977 - loss: 0.4380 - val_accuracy: 0.8403 - val_loss: 0.3749
Epoch 3/20

```

```

469/469          3s 6ms/step -
accuracy: 0.8413 - loss: 0.3583 - val_accuracy: 0.8266 - val_loss: 0.4148
Epoch 4/20

```

```

469/469          3s 6ms/step -
accuracy: 0.8680 - loss: 0.3078 - val_accuracy: 0.7900 - val_loss: 0.5438
Epoch 5/20

```

```

469/469          3s 6ms/step -
accuracy: 0.8901 - loss: 0.2679 - val_accuracy: 0.8493 - val_loss: 0.4226
Epoch 6/20

```

```

469/469          3s 6ms/step -
accuracy: 0.9076 - loss: 0.2313 - val_accuracy: 0.8478 - val_loss: 0.4241
Epoch 7/20

```

```

469/469          3s 6ms/step -
accuracy: 0.9264 - loss: 0.1944 - val_accuracy: 0.8377 - val_loss: 0.5340
Epoch 8/20

```

```

469/469          3s 6ms/step -

```

```

accuracy: 0.9378 - loss: 0.1577 - val_accuracy: 0.8444 - val_loss: 0.5812
Epoch 9/20
469/469          3s 6ms/step -
accuracy: 0.9556 - loss: 0.1160 - val_accuracy: 0.8325 - val_loss: 0.6704
Epoch 10/20
469/469          3s 6ms/step -
accuracy: 0.9687 - loss: 0.0865 - val_accuracy: 0.8378 - val_loss: 0.7235
Epoch 11/20
469/469          3s 6ms/step -
accuracy: 0.9794 - loss: 0.0621 - val_accuracy: 0.8382 - val_loss: 0.8978
Epoch 12/20
469/469          3s 6ms/step -
accuracy: 0.9837 - loss: 0.0465 - val_accuracy: 0.8224 - val_loss: 0.8137
Epoch 13/20
469/469          3s 6ms/step -
accuracy: 0.9857 - loss: 0.0392 - val_accuracy: 0.8327 - val_loss: 0.9521
Epoch 14/20
469/469          3s 5ms/step -
accuracy: 0.9930 - loss: 0.0253 - val_accuracy: 0.8318 - val_loss: 1.1659
Epoch 15/20
469/469          3s 5ms/step -
accuracy: 0.9925 - loss: 0.0249 - val_accuracy: 0.8299 - val_loss: 1.1141
Epoch 16/20
469/469          3s 6ms/step -
accuracy: 0.9941 - loss: 0.0177 - val_accuracy: 0.8146 - val_loss: 1.5620
Epoch 17/20
469/469          2s 5ms/step -
accuracy: 0.9944 - loss: 0.0166 - val_accuracy: 0.8309 - val_loss: 1.1767
Epoch 18/20
469/469          2s 5ms/step -
accuracy: 0.9943 - loss: 0.0133 - val_accuracy: 0.8320 - val_loss: 1.2911
Epoch 19/20
469/469          3s 5ms/step -
accuracy: 0.9948 - loss: 0.0169 - val_accuracy: 0.8114 - val_loss: 1.4533
Epoch 20/20
469/469          3s 5ms/step -
accuracy: 0.9955 - loss: 0.0143 - val_accuracy: 0.8360 - val_loss: 1.5275

```

```
[ ]: <keras.src.callbacks.history.History at 0x7ada181c6990>
```

Testing full transformer

```
[ ]: model = keras.models.load_model(
    "full_transformer_encoder02.keras",
    custom_objects={"TransformerEncoder": TransformerEncoder,
                    "PositionalEmbedding": PositionalEmbedding})
print(f"Test acc: {model.evaluate(int_test_ds)[1]:.3f}")
```

```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/layer.py:393:
UserWarning: `build()` was called on layer 'positional_embedding', however the
layer does not have a `build()` method implemented and it looks like it has
unbuilt state. This will cause the layer to be marked as built, despite not
being actually built, which may cause failures down the line. Make sure to
implement a proper `build()` method.
  warnings.warn(
/usr/local/lib/python3.11/dist-packages/keras/src/layers/layer.py:393:
UserWarning: `build()` was called on layer 'transformer_encoder_1', however the
layer does not have a `build()` method implemented and it looks like it has
unbuilt state. This will cause the layer to be marked as built, despite not
being actually built, which may cause failures down the line. Make sure to
implement a proper `build()` method.
  warnings.warn(

782/782          6s 6ms/step -
accuracy: 0.8326 - loss: 0.3761
Test acc: 0.830

```

Transformer with pretrained embedding

```

[ ]: #Vectorizing the data

max_length = 150
max_tokens = 10000
text_vectorization = layers.TextVectorization(
    max_tokens=max_tokens,
    output_mode="int",
    output_sequence_length=max_length,
)
text_vectorization.adapt(text_only_train_ds)

int_train_ds = train_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
int_val_ds = val_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
int_test_ds = test_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
#Transformer encoder implemented as a subclassed Layer

class TransformerEncoder(layers.Layer):
    def __init__(self, embed_dim, dense_dim, num_heads, **kwargs):
        super().__init__(**kwargs)
        self.embed_dim = embed_dim
        self.dense_dim = dense_dim

```

```

        self.num_heads = num_heads
        self.attention = layers.MultiHeadAttention(
            num_heads=num_heads, key_dim=embed_dim)
        self.dense_proj = keras.Sequential(
            [layers.Dense(dense_dim, activation="relu"),
             layers.Dense(embed_dim),]
        )
        self.layernorm_1 = layers.LayerNormalization()
        self.layernorm_2 = layers.LayerNormalization()

    def call(self, inputs, mask=None):
        if mask is not None:
            mask = mask[:, tf.newaxis, :]
        attention_output = self.attention(
            inputs, inputs, attention_mask=mask)
        proj_input = self.layernorm_1(inputs + attention_output)
        proj_output = self.dense_proj(proj_input)
        return self.layernorm_2(proj_input + proj_output)

    def get_config(self):
        config = super().get_config()
        config.update({
            "embed_dim": self.embed_dim,
            "num_heads": self.num_heads,
            "dense_dim": self.dense_dim,
        })
        return config
#Using the Transformer encoder for text classification

vocab_size = 10000
embed_dim = 100
num_heads = 2
dense_dim = 32

inputs = keras.Input(shape=(None,), dtype="int64")
x = embedding_layer(inputs)
#x = layers.Embedding(vocab_size, embed_dim)(inputs)
x = TransformerEncoder(embed_dim, dense_dim, num_heads)(x)
x = layers.GlobalMaxPooling1D()(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
model.summary()
#Training and evaluating the Transformer encoder based model

```

```
callbacks = [
    keras.callbacks.ModelCheckpoint("transformer_encoder_pretrained{epoch:02d}.
↳keras",
                                   save_best_only=True)
]
model.fit(int_train_ds, validation_data=int_val_ds, epochs=20,↳
↳callbacks=callbacks)
```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/layer.py:938:  
 UserWarning: Layer 'transformer\_encoder\_3' (of type TransformerEncoder) was  
 passed an input with a mask attached to it. However, this layer does not support  
 masking and will therefore destroy the mask information. Downstream layers will  
 not see the mask.

```
warnings.warn(
```

Model: "functional\_15"

Layer (type)	Output Shape	Param #	Connected to
input_layer_15 (InputLayer)	(None, None)	0	-
embedding_3 (Embedding)	(None, None, 100)	1,000,000	input_layer_15[0...
not_equal_7 (NotEqual)	(None, None)	0	input_layer_15[0...
transformer_encode... (TransformerEncode...	(None, None, 100)	87,632	embedding_3[2][0... not_equal_7[0][0]
global_max_pooling... (GlobalMaxPooling1...	(None, 100)	0	transformer_enco...
dropout_15 (Dropout)	(None, 100)	0	global_max_pooli...
dense_24 (Dense)	(None, 1)	101	dropout_15[0][0]

Total params: 1,087,733 (4.15 MB)

Trainable params: 87,733 (342.71 KB)

Non-trainable params: 1,000,000 (3.81 MB)

Epoch 1/20

469/469 30s 37ms/step -

accuracy: 0.5931 - loss: 0.7832 - val\_accuracy: 0.7715 - val\_loss: 0.4730

Epoch 2/20

469/469 4s 8ms/step -

accuracy: 0.7905 - loss: 0.4596 - val\_accuracy: 0.8086 - val\_loss: 0.4129

Epoch 3/20

469/469 4s 8ms/step -

accuracy: 0.7987 - loss: 0.4379 - val\_accuracy: 0.8150 - val\_loss: 0.4082

Epoch 4/20

469/469 2s 4ms/step -

accuracy: 0.8067 - loss: 0.4272 - val\_accuracy: 0.8011 - val\_loss: 0.4234

Epoch 5/20

469/469 4s 8ms/step -

accuracy: 0.8108 - loss: 0.4162 - val\_accuracy: 0.8236 - val\_loss: 0.3912

Epoch 6/20

469/469 2s 5ms/step -

accuracy: 0.8161 - loss: 0.4112 - val\_accuracy: 0.8226 - val\_loss: 0.3984

Epoch 7/20

469/469 2s 4ms/step -

accuracy: 0.8198 - loss: 0.4045 - val\_accuracy: 0.8149 - val\_loss: 0.4058

Epoch 8/20

469/469 4s 9ms/step -

accuracy: 0.8199 - loss: 0.4025 - val\_accuracy: 0.8211 - val\_loss: 0.3901

Epoch 9/20

469/469 2s 4ms/step -

accuracy: 0.8229 - loss: 0.3954 - val\_accuracy: 0.8161 - val\_loss: 0.4064

Epoch 10/20

469/469 2s 4ms/step -

accuracy: 0.8281 - loss: 0.3888 - val\_accuracy: 0.8184 - val\_loss: 0.3979

Epoch 11/20

469/469 2s 4ms/step -

accuracy: 0.8334 - loss: 0.3820 - val\_accuracy: 0.8164 - val\_loss: 0.3982

Epoch 12/20

469/469 2s 5ms/step -

accuracy: 0.8335 - loss: 0.3762 - val\_accuracy: 0.8168 - val\_loss: 0.3978

Epoch 13/20

469/469 2s 4ms/step -

accuracy: 0.8358 - loss: 0.3713 - val\_accuracy: 0.8058 - val\_loss: 0.4348

Epoch 14/20

469/469 2s 4ms/step -

accuracy: 0.8422 - loss: 0.3634 - val\_accuracy: 0.8209 - val\_loss: 0.3988

Epoch 15/20

469/469 2s 4ms/step -

accuracy: 0.8414 - loss: 0.3596 - val\_accuracy: 0.8008 - val\_loss: 0.4477

```

Epoch 16/20
469/469          2s 4ms/step -
accuracy: 0.8465 - loss: 0.3578 - val_accuracy: 0.8143 - val_loss: 0.4074
Epoch 17/20
469/469          2s 4ms/step -
accuracy: 0.8473 - loss: 0.3485 - val_accuracy: 0.8148 - val_loss: 0.4166
Epoch 18/20
469/469          2s 4ms/step -
accuracy: 0.8512 - loss: 0.3418 - val_accuracy: 0.8066 - val_loss: 0.4387
Epoch 19/20
469/469          2s 4ms/step -
accuracy: 0.8578 - loss: 0.3361 - val_accuracy: 0.7989 - val_loss: 0.4434
Epoch 20/20
469/469          2s 4ms/step -
accuracy: 0.8570 - loss: 0.3327 - val_accuracy: 0.8068 - val_loss: 0.4399

```

```
[ ]: <keras.src.callbacks.history.History at 0x7ada28702d50>
```

Testing transformer with pretrained embedding

```
[ ]: model = keras.models.load_model(
    "transformer_encoder_pretrained08.keras",
    custom_objects={"TransformerEncoder": TransformerEncoder})
print(f"Test acc: {model.evaluate(int_test_ds)[1]:.3f}")
```

```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/layer.py:393:
UserWarning: `build()` was called on layer 'transformer_encoder_3', however the
layer does not have a `build()` method implemented and it looks like it has
unbuilt state. This will cause the layer to be marked as built, despite not
being actually built, which may cause failures down the line. Make sure to
implement a proper `build()` method.
  warnings.warn(

```

```

782/782          5s 5ms/step -
accuracy: 0.8148 - loss: 0.3956
Test acc: 0.813

```

Full transformer with pretrained embedding

```
[ ]: #Using positional encoding to re-inject order information
#Implementing positional embedding as a subclassed layer

class PositionalEmbedding(layers.Layer):
    def __init__(self, sequence_length, input_dim, output_dim, **kwargs):
        super().__init__(**kwargs)
        self.token_embeddings = layers.Embedding(
            input_dim=input_dim,
            output_dim=output_dim,
```



```

        embeddings_initializer=keras.initializers.
↪Constant(embedding_matrix),
        trainable=False,
        mask_zero=True,
    )
    self.position_embeddings = layers.Embedding(
        input_dim=sequence_length, output_dim=output_dim)
    self.sequence_length = sequence_length
    self.input_dim = input_dim
    self.output_dim = output_dim

    def call(self, inputs):
        length = tf.shape(inputs)[-1]
        positions = tf.range(start=0, limit=length, delta=1)
        embedded_tokens = self.token_embeddings(inputs)
        embedded_positions = self.position_embeddings(positions)
        return embedded_tokens + embedded_positions

    def compute_mask(self, inputs, mask=None):
        return keras.ops.not_equal(inputs, 0)

    def get_config(self):
        config = super().get_config()
        config.update({
            "output_dim": self.output_dim,
            "sequence_length": self.sequence_length,
            "input_dim": self.input_dim,
        })
        return config

#Putting it all together: A text-classification Transformer
#Combining the Transformer encoder with positional embedding

vocab_size = 10000
sequence_length = 150
embed_dim = 100
num_heads = 2
dense_dim = 32

inputs = keras.Input(shape=(None,), dtype="int64")
x = PositionalEmbedding(sequence_length, vocab_size, embed_dim)(inputs)
x = TransformerEncoder(embed_dim, dense_dim, num_heads)(x)
x = layers.GlobalMaxPooling1D()(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",

```

```

        metrics=["accuracy"])
model.summary()

callbacks = [
    keras.callbacks.ModelCheckpoint("full_transformer_encoder_pretrained{epoch:
    ↪02d}.keras",
                                   save_best_only=True)
]
model.fit(int_train_ds, validation_data=int_val_ds, epochs=20,
    ↪callbacks=callbacks)

```

Testing full transformer with pretrained embedding

```

[ ]: model = keras.models.load_model(
    "full_transformer_encoder_pretrained15.keras",
    custom_objects={"TransformerEncoder": TransformerEncoder,
                    "PositionalEmbedding": PositionalEmbedding})
print(f"Test acc: {model.evaluate(int_test_ds)[1]:.3f}")

```